

Relational Database Design - Lecture 5

Steps : ER \rightarrow Schema

1. Regular Entity Types Mapping

2. Weak Entity Types

- Create table separate table for each regular entity types.
- Simple Attribute \rightarrow Column
- Simple Components of composite \rightarrow Column
- Multi-Valued \rightarrow Refer to step 6
- Derived \rightarrow Dont Add
- PK \rightarrow (1 candidate key) will be selected and underlined
or for composite attributes, underline all simple components

2. Weak Entity Type Mapping

- Follow all from step 1 except for PK.
- Owner's PK \rightarrow Weak Entity's FK (shown by arrows)
If this was composite, then add all as FK.
- \rightarrow PK \rightarrow FK + Weak Entity's Partial key
 - Underline both

3. Binary N:M Relationship Types Mapping

- Create a separate table for each N:M binary relationship
- PK of Connected Entities \rightarrow FK in relationship table
- PK \rightarrow The combined FK, and should be underlined
- Add Simple and Composite field of relationship like step 1.

Cross Referencing : Method of creating separate table for relationships

4. Binary 1:N Relationship Mapping

- Cross Referencing is inefficient, use foreign key approach.
- For every 1:N relationship
 - PK of table in 1's side \rightarrow FK of table in N's side
- Add relationship attributes to N side of relationship.

5. 1:1 Relationship Mapping

- 3 Possible Approach:

- Cross Referencing: Same step as step 3

- Merged Relation Option:

- Merge two entity tables into one

- PK: Combination of two PK

- May lead to data redundancy

- Foreign Key Approach: (Best Approach)

- take any of the table's PK and add as FK for other's table

- Add as FK to the table which has a "tot" participation, to prevent null values.

Recursive Relationship

- Use same approach as binary 1:N, 1:1, N:M.

step 4

step 5

step 3

6. Multivalued Attribute Mapping

- Each Multivalued Attribute \rightarrow Separate table (including multivalued composite)
- PK of entity table \rightarrow FK of new table
- Multivalued Only:
 - Add another column for storing value
 - Combined Primary key: FK and the column in prev line.

- Multivalued Composite:

- Add only simple components of attribute as separate column
- Combined PK: Fk + Partial Key (if exists)
- If no PK: Underline all components as part of PK

7. N-ary Relationship Mappings

- Use cross-referencing method like step 3.

8. Mapping Specialization or Generalization

→ Convert each specialization

Consider the subclass S₁, S₂, S₃ and superclass C.

- (a) Multiple relations - superclass and subclass : Works for all type
 - C - Has the PK and all common fields
 - S₁, S₂, S₃ - Has the PK and the uncommon fields
- (b) Multiple relations - subclass relations only : Only works for total
 - C - No table
 - S₁, S₂, S₃ - Separate tables with common fields and uncommon fields
- (c) Single relation with one type attribute : Disjoint
 - C - Has all the fields (some will be null) of S₁, S₂, S₃ and one ~~attr~~ attribute called a type [to denote ~~s1/s2/s3~~]
- (d) Single relation with multiple type attributes : Overlap
 - C - Will have all fields of S₁, S₂, S₃. There will be S₁Flag, S₂Flag, S₃Flag to denote overlap.
- (e) Mapping of shared subclass (Multiple inheritance)
 - You can use any of the options above (Note: The superclass may have the same PK)

Normalization - Lecture 6

A process used to make database more efficient.

Informal Guideline - No written rules.

Formal Guidelines

Normal Forms

- 1st NF, 2nd NF, ~~3rd NF~~, BCNF, NF, 3rd NF, 4th NF, 5th NF

1st Normal Form

- Disallows
 - Composite Attribute
 - Multivalued Attribute
 - Nested Relation; Attributes whose values for an individual tuple are non-atomic
- Considered to be a part of the definition of relation.

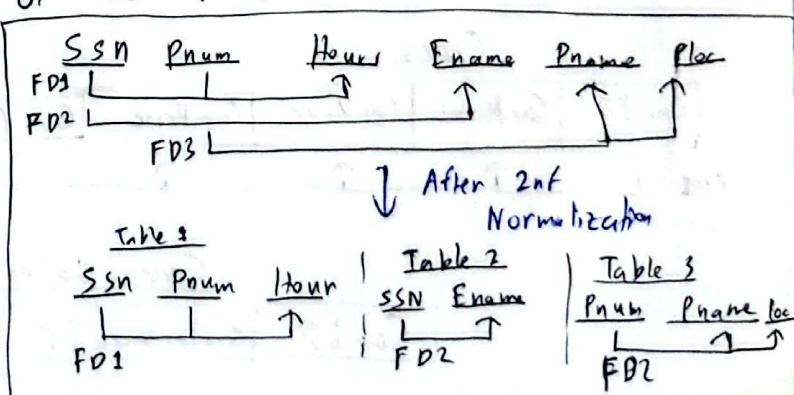
2nd Normal Form

- **Prime Attribute**: An attribute that is member of ~~part of primary~~ candidate key K.
- **Full Functional Dependency**: A FD (functional dependency) $Y \rightarrow Z$ where removal of any attribute from Y means the FD does not hold anymore.
- One or more attributes depend fully on the PK

Partial Functional Dependency: One or more attributes depend partly on ~~the PK~~ a part of the PK.

Candidate key: Every key with potential to be PK

Secondary key: Candidate key not chosen as PK



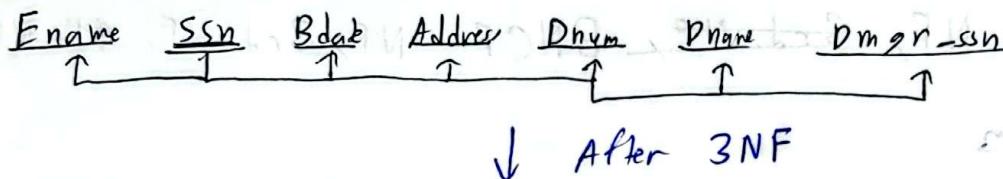
3rd Normal Form

Transitive FD: If $X \rightarrow Y$, $Y \rightarrow Z$ then $X \rightarrow Z$

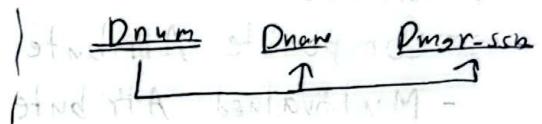
Condition

- It must be in 2NF
- Then ~~must~~ cannot be any transitive FD.

Example

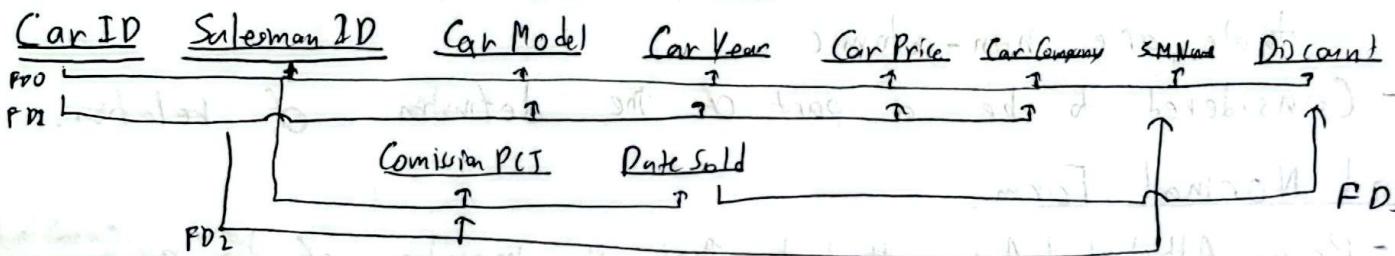


Ename	Ssn	Bdate	Address	Dnum
↑	↑	↑	↑	↑



Question Eg

Car Sales Table



Q) Is this in 1st NF?

- Yes. There are no multivalued, composite attributes or nested relations.

Q) Is this in 2nd NF? Why or why not.

Convert to 3rd NF.

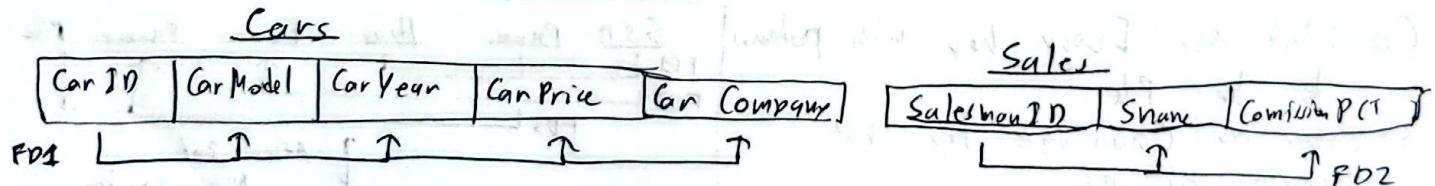
It is not in 2 NF as FD1, FD2 has partial dependency.

Some Notations

Car Model (Name, Year)

Composite

{Car Model} - Multivalued



Car-Sales

Car ID	Salesman ID	Discount	Date Sold
↑	↑	↑	↑

FDD: Car ID → Salesman ID, Salesman ID → Discount, Salesman ID → Date Sold

- Q) Explain if the relational schema is in 3NF or not. If not, normalize to 3NF.
- It is not in 3NF as there is transitive dependency, FD₃.

Car

CarID	Car Model	Car Year	Car Price	Car Company
FD ₁		↑	↑	↑

Sales

SalesmanID	SMName	Commission Pct
FD ₂		↑

CarSales

CarID	SalesmanID	Date Sold
FD ₂	↑	↑

Discount

Date Sold	Discount
↑	↑ FD ₃

Note: BCNF is not in syllabus

Superkey: Set of attribute which will be unique across all rows

- Any ~~key~~ K is a superkey if removing any attribute from K will cause K not to be a superkey any more.

Indexing in Database - Lecture 7

- An Index in database works in same way as textbook
- Speeds up access to desired data

Search key: Attribute or set of attributes used to look up records in a file.

Index file:

Search-key	pointer
------------	---------

- Index files are smaller than original file

Types of indices:

- Ordered indices: Clustering Index, Non clustering Index
- Hash indices: Search keys are distributed uniformly across "buckets" using a "hash function"

Ordered Indices

- Index entries are stored sorted on the search key value

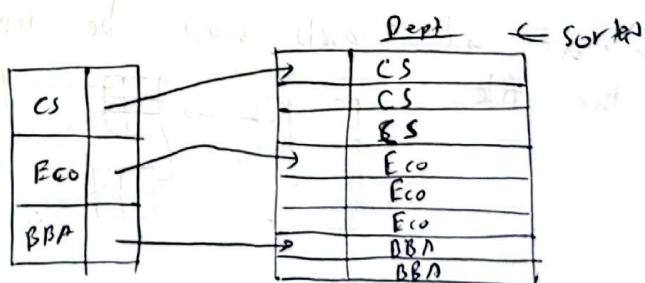
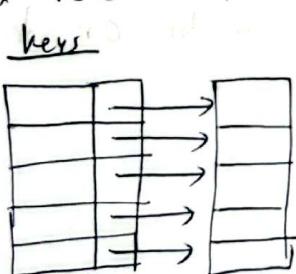
Clustering Index: In a sequentially ordered file, the index whose search key specifies the sequential order of the file

Secondary Index: An index whose search key specifies an order different from the sequential order of the file.

Index-sequential file: Sequential file ordered on a search key with a clustering index on the search key.

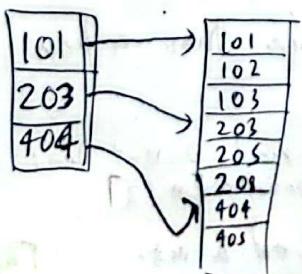
Dense Index Files

Index record appears for every search-key value in the file



Sparse Index Files

- Contains index records for only some search-key values
- Applicable when records are stored sequentially



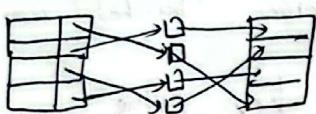
- Searching for k :
 - Find largest key $< k$
 - Sequentially continue

- Less space and less maintenance overhead for insertion and deletion
- Generally slower than dense index for locating records.
- Good Tradeoff:

For Clustered Index: Sparse index with an index entry for every block in file corresponding to least search-key value in the block



For Unclustered Index: Sparse index on top of dense index (multilevel index)

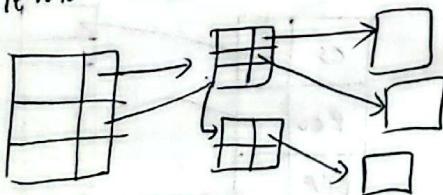


Secondary Indices

- Index record points to a bucket that contains pointers to all the actual records with that particular search-key value
- Secondary indices have to be dense

Multilevel Index

- If index does not fit in memory, access becomes expensive
- Solution: Treat index kept on disk as a sequential file and construct a sparse index on it.
- Outer Index - A sparse index of the basic index
- Inner Index - The basic index file
- Even if outer index is too large to fit in main memory yet another level of index can be created, and so on.
- Indices at all levels must be updated on insertion or deletion from the file.



B+ Tree

- Every tree has an order of "n"
- All nodes on the left have smaller values than a key and all nodes on the right have greater than or equal values than a key.

Each node has

- Children
- Keys/Values

Max	Min
• root/Internal : n leaf : 0	• Root : 2 (if root is not leaf as well) Internal : $\lceil \frac{n}{2} \rceil$
• n-1	• Root : 1, Leaf & Internal : $\lceil \frac{n}{2} \rceil - 1$

Note: for simplicity leaf and internal is

considered same in some resources

otherwise leaf min $\Omega(n-1)/2\Gamma$

B+ Tree (Insertion)

- 1) Find the right spot (leaf node): Start from root and move down the tree, following the correct path (based on the key you want to insert), until you reach a leaf.
- 2) If node has less than "max" keys:
 - Insert value in sorted order
- 3) If node has "max" keys:
 - Split into half
 - 1st half: New left leaf
 - 2nd half: New right leaf
 - Parent: lowest key from right node
- 4) Repeat split upward if needed.
 - If parent now becomes full too, split it as well and push a key further up.
 - If a non-leaf node is split, the lowest value on the right will only be moved up and not copied to the new right leaf. This might continue up to the root.
 - If root splits \rightarrow new root will be created and the tree height will increase.
- If order of B+ tree is odd, then create a right-biased or left-biased tree. Maintain the same bias for entire tree.

B+ Tree (Deletion)

- Use common sense and make sure the properties of B+ tree are maintained

B+ Properties

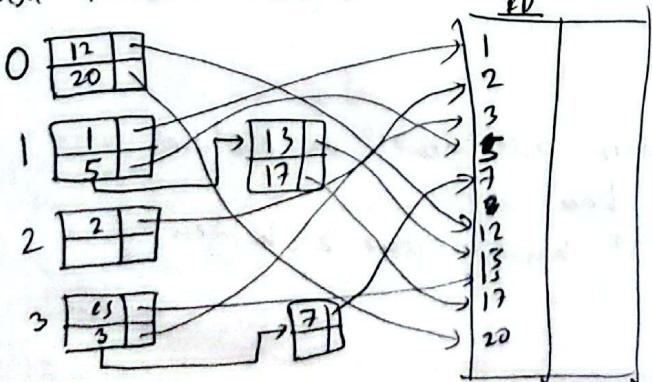
- Balanced: All leaf at same level
- Order m: Each internal node has at most m children
- Leaf stores at most $m-1$ keys
- If root is not leaf, it has atleast 2 children

Static Hashing

- Bucket: Unit of storage containing one or more entries (data block)
 - Obtained from an entry's search-key value using a hash function
- Hash Function, h : Set of all search-key values $K \rightarrow$ Set of all bucket addresses B
 - Used to locate entries for access, insertion as well as deletion
 - Entries with different search-key values may be mapped to the same bucket, thus entire bucket has to be searched sequentially to locate an entry.
 - hash index : buckets store entries with pointers to records
 - hash file organization : buckets ^{buckets overflow} store records
 - Bucket overflow can occur because of
 - Insufficient buckets
 - Skew in distribution records. This can occur due to two reasons:
 - Multiple records have same search-key value
 - Choose hash function produces non-uniform distribution of key values
 - Although the probability of bucket overflow can be reduced, it cannot be eliminated; it is handled by using overflow buckets
 - Overflow Chaining - The overflow buckets of a given bucket are chained together in a linked list
 - Above scheme is called closed addressing
 - An alternative, called open addressing which does not use overflow buckets, is not suitable for database applications.

Example

Hash Function $h(ID) = (\text{sum of digits in ID}) \% 4$



Transaction, Concurrency Control & Recovery

Transaction: Logical unit of database processing that includes one or more access operations

BEGIN TRANSACTION;

UPDATE accounts SET balance = balance - 50 WHERE account-id = 'A';

and simultaneously another transaction + 50 " " = 'B';

COMMIT

- Transaction may be stand-alone or may be embedded within a program

ACID

• Used to preserve the integrity of data enforced by Concurrency Control and recovery methods.

Atomicity: A transaction is an atomic unit of processing; it should either be performed in its entirety or not performed at all. Ensured by recovery, recovery subsystem.

Consistency: A transaction should be consistency preserving, meaning if it is completely executed from beginning to end, it should take the database from one, consistent state to another.

Isolation: Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrent transactions executing transactions. A transaction should not make its update visible to other transactions until it is committed. Ensured by running transaction serially (concurrent control).

Durability: The changes applied to the database by a committed transaction must persist in the database. These changes must not be lost because of any failure. Ensured by recovery subsystem of DBMS.

Issues that can occur for Transaction

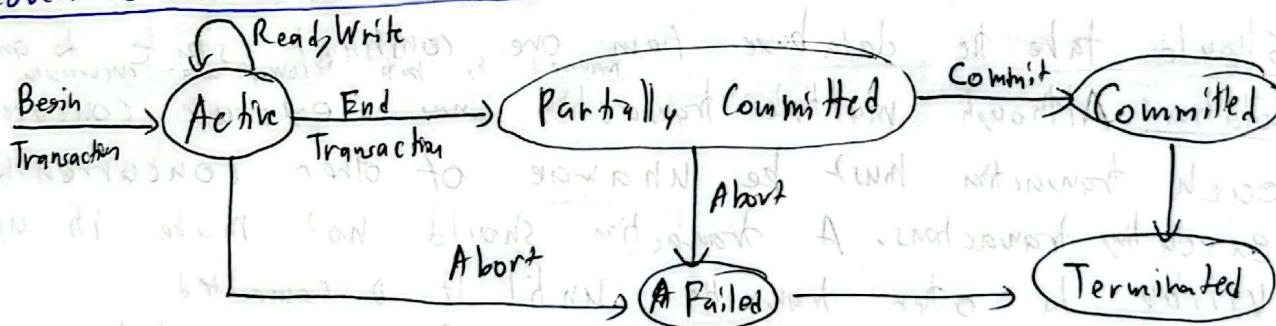
- Recovery: Take DB to previous stable state

- Concurrency Control: Concurrent execution of multiple transaction

Why Recovery is Needed

- Computer Failure (system crash)
- A transaction or system error: Logical error, erroneous params, interrupt
 - Local errors or exception detected by the transaction
 - Certain conditions necessitate cancellation of transaction
 - [Data of transaction not found ; insufficient balance
 - [A to programmed abort in the transaction.
- Concurrency control enforcement
 - Concurrency control method may decide to abort the transaction, to be restarted later, because it violates serializability or because several transactions are in the state of deadlock
- Disk Failure
 - Physical Problem and Catastrophes: Includes power, fire, etc.

Recovery (Transaction States)



Concurrency Control

- Multiple transaction are allowed to run concurrently in the system
- Advantages:
 - Increased processor and disk utilization
 - Reduced average response time for transaction

Concurrency Control Schemes: Mechanisms to achieve isolation to prevent transaction from destroying consistency of database.

Issues that might be caused in absence of Concurrency Control

- **Lost Update Problem:** Two transaction updating same data, cause one getting overwritten.
- **Dirty Read (Temporary Update) Problem:** Transaction reads ~~uncommitted~~ data written by uncommitted transaction and the ~~transac~~ uncommitted transaction gets reverted.
- **The Incorrect Summary (or Incorrect Analysis) Problem:** A transaction performs an aggregate function while other transactions are updating data. The aggregate function may calculate some values before they are updated and others after they are updated.
- **Non Repeatable Read:** A transaction reads the same row twice and get different value due to another transaction.
- **Phantom Read:** A transaction reexecuting a query ~~but~~ ~~getting~~ returning a set of rows, rows appear or disappear due to another transaction's insert or deletes.

Transaction in SQL

- SQL statements are always atomic
- End statements: COMMIT, ROLLBACK
 - By default. Can be turned off.
- Isolation level can be changed at the start of transaction or be set at database level.

Isolation Levels

<u>Isolation Level</u>	<u>Dirty Read</u>	<u>Non-Repeatable Read</u>	<u>Phantom Read</u>
<u>Serializable</u> : Transactions happen in a serial order.	Prevents	Prevents	Prevents
<u>Repeatable Read</u> : Only committed records can be read. Repeated read, allows ^{same value}	Prevents	Prevents	Allows
<u>Read Committed</u> : Only committed records can be read.	Prevents	Allows	Allows
<u>Read Uncommitted</u> : Even uncommitted records can be read.	Allows	Allows	Allows